

1.

We are given as input a set of n requests (e.g., for the use of an auditorium), with a known start time s_i and finish time t_i for each request i . Assume that all start and finish times are distinct. Two requests *conflict* if they overlap in time --- if one of them starts between the start and finish times of the other. Our goal is to select a maximum-cardinality subset of the given requests that contains no conflicts. (For example, given three requests consuming the intervals $[0, 3]$, $[2, 5]$, and $[4, 7]$, we want to return the first and third requests.) We aim to design a greedy algorithm for this problem with the following form: At each iteration we select a new request i , including it in the solution-so-far and deleting from future consideration all requests that conflict with i .

Which of the following greedy rules is guaranteed to always compute an optimal solution?

☐

At each iteration, pick the remaining request which requires the least time (i.e., has the smallest value of $t_i - s_i$) (breaking ties arbitrarily).

☐

At each iteration, pick the remaining request with the fewest number of conflicts with other remaining requests (breaking ties arbitrarily).

☐

At each iteration, pick the remaining request with the earliest start time.

☒

At each iteration, pick the remaining request with the earliest finish time.

正确

Let R_j denote the requests with the j earliest finish times. Prove by induction on j that this greedy algorithm selects the maximum-number of non-conflicting requests from S_j .

2.

We are given as input a set of n jobs, where job j has a processing time p_j and a deadline d_j . Recall the definition of *completion times* C_j from the video lectures. Given a schedule (i.e., an ordering of the jobs), we define the *lateness* l_j of job j as the amount of time $C_j - d_j$ after its deadline that the job completes, or as 0 if $C_j \leq d_j$. Our goal is to minimize the maximum lateness, $\max_j l_j$.

Which of the following greedy rules produces an ordering that minimizes the maximum lateness? You can assume that all processing times and deadlines are distinct.

☐

None of the other answers are correct.

☒

Schedule the requests in increasing order of deadline d_j

正确

Proof by an exchange argument, analogous to minimizing the weighted sum of completion times.

☐

Schedule the requests in increasing order of the product $d_j \cdot p_j$

☐

Schedule the requests in increasing order of processing time p_j

3.

In this problem you are given as input a graph $T = (V, E)$ that is a tree (that is, T is undirected, connected, and acyclic). A *perfect matching* of T is a subset $F \subset E$ of edges such that every vertex $v \in V$ is the endpoint of exactly one edge of F . Equivalently, F matches each vertex of T with exactly one other vertex of T . For example, a path graph has a perfect matching if and only if it has an even number of vertices.

Consider the following two algorithms that attempt to decide whether or not a given tree has a perfect matching. The *degree* of a vertex in a graph is the number of edges incident to it. (The two algorithms differ only in the choice of v in line 5.)

Algorithm A:

```
1 While T has at least one vertex:
2   If T has no edges:
3     halt and output "T has no perfect matching."
4   Else:
5     Let v be a vertex of T with maximum degree.
6     Choose an arbitrary edge e incident to v.
7     Delete e and its two endpoints from T.
8 [end of while loop]
9 Halt and output "T has a perfect matching."
```

Algorithm B:

```
1 While T has at least one vertex:
2   If T has no edges:
3     halt and output "T has no perfect matching."
4   Else:
5     Let v be a vertex of T with minimum non-zero degree.
6     Choose an arbitrary edge e incident to v.
7     Delete e and its two endpoints from T.
8 [end of while loop]
9 Halt and output "T has a perfect matching."
```

Is either algorithm correct?

☐

Algorithm A always correctly determines whether or not a given tree graph has a perfect matching; algorithm B does not.

☐

Both algorithms always correctly determine whether or not a given tree graph has a perfect matching.

☒

Algorithm B always correctly determines whether or not a given tree graph has a perfect matching; algorithm A does not.

正确

Algorithm A can fail, for example, on a three-hop path. Correctness of algorithm B can be proved by induction on the number of vertices in T . Note that the tree property is used to argue that there must be a vertex with degree 1; if there is a perfect matching, it must include the edge incident to this vertex.

☐

Neither algorithm always correctly determines whether or not a given tree graph has a perfect matching.

4.

Consider an undirected graph $G = (V, E)$ where every edge $e \in E$ has a given cost c_e . Assume that all edge costs are positive and distinct. Let T be a minimum spanning tree of G and P a shortest path from the vertex s to the vertex t . Now suppose that the cost of every edge e of G is increased by 1 and becomes $c_e + 1$. Call this new graph G' . Which of the following is true about G' ?

☐

T may not be a minimum spanning tree and P may not be a shortest s - t path.

☐

T may not be a minimum spanning tree but P is always a shortest s - t path.

☐

T is always a minimum spanning tree and P is always a shortest s - t path.

☒

T must be a minimum spanning tree but P may not be a shortest s - t path.

正确

The positive statement has many proofs (e.g., via the Cut Property). For the negative statement, think about two different paths from s to t that contain a different number of edges.

5.

Suppose T is a minimum spanning tree of the connected graph G . Let H be a connected induced subgraph of G . (i.e., H is obtained from G by taking some subset $S \subseteq V$ of vertices, and taking all edges of E that have both endpoints in S . Also, assume H is connected.) Which of the following is true about the edges of T that lie in H ? You can assume that edge costs are distinct, if you wish. [Choose the strongest true statement.]

☒

For every G and H , these edges are contained in some minimum spanning tree of H

正确

Proof via the Cut Property (cuts in G correspond to cuts in H with only fewer crossing edges).

☐

For every G and H , these edges form a minimum spanning tree of H

☐

For every G and H and spanning tree T_H of H , at least one of these edges is missing from T_H

☐

For every G and H , these edges form a spanning tree (but not necessary minimum-cost) of H